

Neural Networks

Jan Kościółkowski



18 May 2020

Idea: approximate some function

- Examples: regression function, conditional probability
- Can be much more complicated
- Define some loss function (discrepancy between true values and predictions) and try to minimize it
- Example: number of claims, Poisson loss, $y = \exp(f(x_1, \dots, x_p))$

Multi-Layer Perceptron (MLP)

- Nodes – apply activation function
- Edges – matrix multiplication
- Universal Approximation Theorem (Cybenko 1989):
Approximates any continuous function on a compact subset of \mathbb{R}^d arbitrarily well




$$f^k = W^k g^{k-1}$$


$$g^k = \sigma(f^k)$$

Activation Functions

Issue: saturation




sigmoid: $\sigma(x) = \frac{1}{1+e^{-x}}$




tanh: $\tanh(x) = \frac{e^x + e^{-x}}{e^x - e^{-x}}$

Issue: dead ReLU

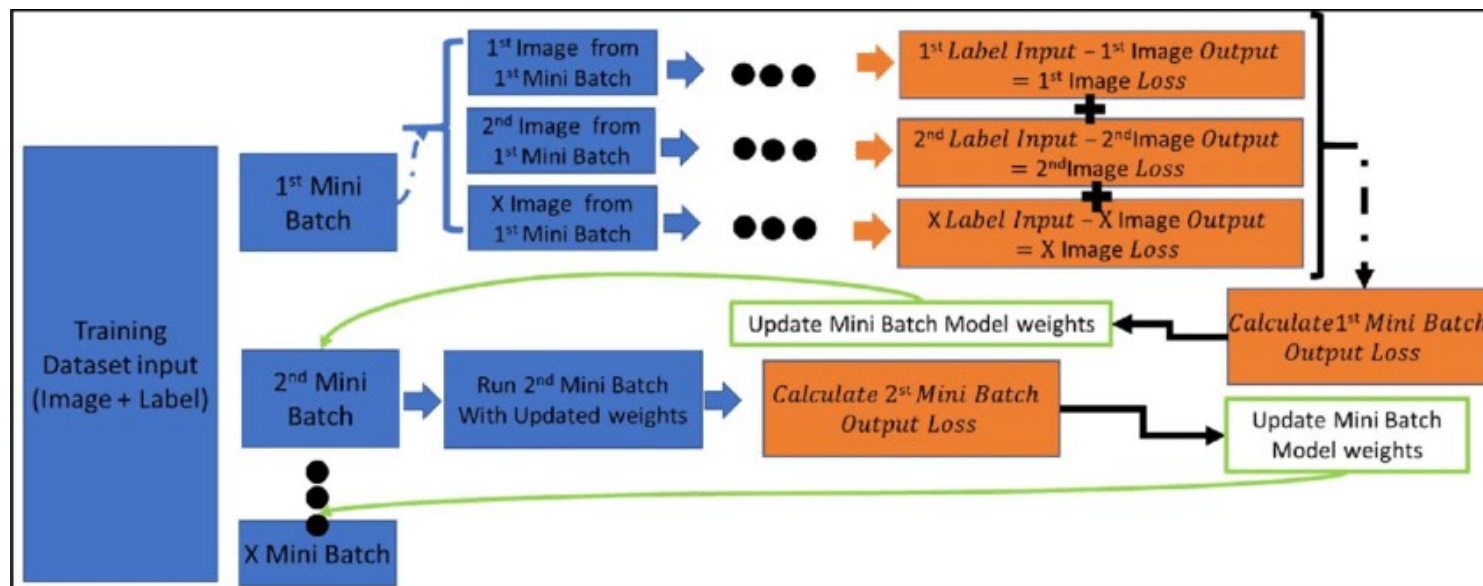


ReLU: $\text{ReLU}(x) = \max(0, x)$



leaky ReLU: $\text{LReLU}(x) = \max(ax, x)$

Mini-batch Learning




https://www.researchgate.net/figure/Example-of-mini-batch-learning-on-CNN_fig41_339054847

Optimization

Neural networks are optimized using (mini-batch) SGD and its variants:

- Momentum – exponential moving average of past gradients
- Learning rate decay – decrease learning rate e.g. according to an exponential schedule
- Adagrad – for each weight w decay learning rate by \sqrt{G} where G is the sum of squares of past gradients for w
- RMSProp – as in Adagrad, but G is exponential moving average of G and squared gradients
- Adam – combines EMA of squared gradients and EMA of gradients, usually the easiest to use

$$w \leftarrow w - \eta \frac{\partial L}{\partial w}$$




<https://rnrahman.com/blog/visualising-stochastic-optimisers/>

Backpropagation

- We need derivatives of loss function w.r.t. weights for SGD
- Use multi-dimensional chain rule to decompose derivative computations into products of step-by-step derivatives
- In general, works for DAGs of differentiable operations (including MLP)
- Idea – for each mini-batch do the following:
 - Forward pass: run the input data through the network and compute all intermediate values
 - Compute the value of the loss function
 - Compute the single-step derivatives going backwards from the loss function
 - Multiply them to get the desired derivatives

$$\frac{\partial L}{\partial z} = \sum_{z=z_0 \rightarrow \dots \rightarrow z_l=L} \prod_{i=1}^{l-1} \frac{\partial z_{i+1}}{\partial z_i}.$$

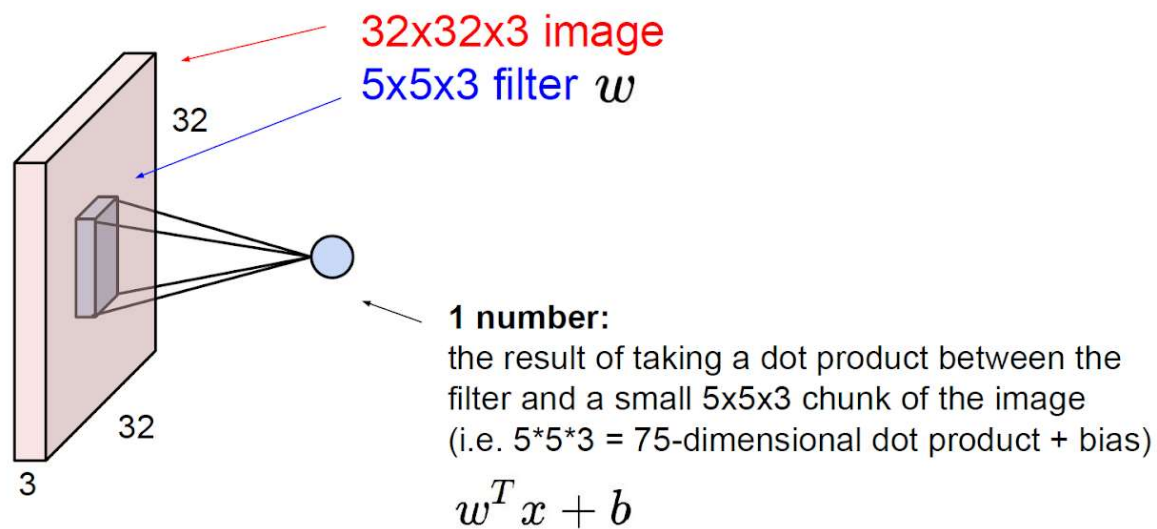


Regularization


NNs are usually too complex for tabular data, heavy regularization needed:

- L1, L2
- Data augmentation – perturb the real data (add noise, offset, distort etc.)
- Dropout – randomly omit network edges during training
- Batch, layer normalization – normalize across different dimensions of the batch to try to make the distributions more normal
- Early stopping – monitor validation set metrics and stop training when they stop improving

Convolutional Networks




<https://cs231n.github.io/convolutional-networks/>




<https://codelabs.developers.google.com/codelabs/cloud-tensorflow-mnist/#10>

Convolutional Networks: History

- Around since 1980s
- LeNet-5 (LeCun et al. 1998) – handwritten digit recognition
- Due to GPU utilization, massive advances since 2012 + switch from SVMs to NNs

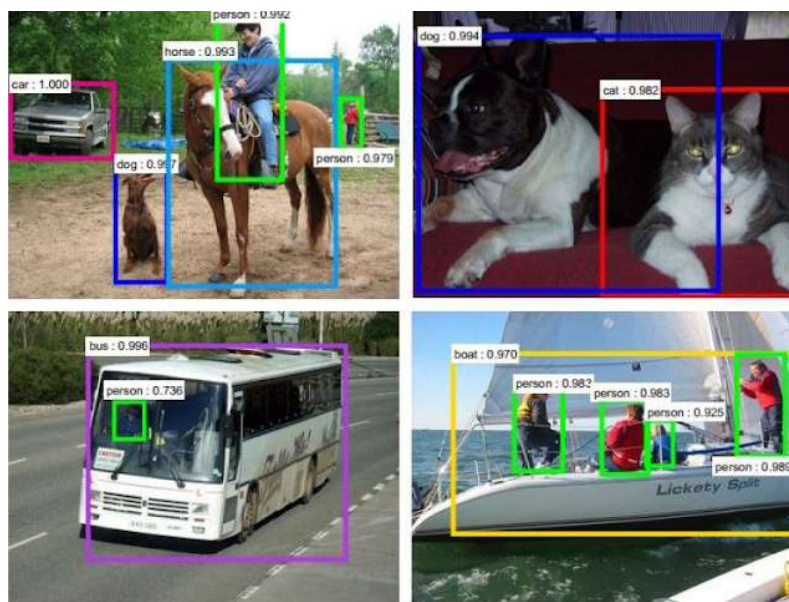


ImageNet Classification Error (Top 5)



Convolutional Networks: Use Cases


Detection



Segmentation



Style Transfer





<https://pjreddie.com/darknet/yolo/>

<https://github.com/xunhuang1995/AdaIN-style>

Recurrent Networks


Used for modeling sequential data: text, voice, time series etc.

Operations exactly the same as for MLP, but uses the concept of weight sharing and backpropagation through time.



Representation Learning

- NNs very good at learning well-performing low-dimensional representations of complex objects (words, images etc.)
- Usually works by trying to reproduce the initial object or the context, the representations are taken to be outputs from some intermediate layer
- Example: autoencoders



<https://towardsdatascience.com/generating-images-with-autoencoders-77fd3a8dd368>

Other Important Classes

- Generative Adversarial Networks (generating fake photos/videos)
- Reinforcement Learning (learning by interaction with environment)
- Transformers (NLP, new models have billions of learnable weights)
- TBC...

Toolkit

Python is the language of choice for Deep Learning research and prototyping, but compiled languages are often used for production uses (e.g. C++, Go)

- Pillow
- NLTK
- PyTorch
 - Pytorch Lightning/Ignite
- Tensorflow (has R API)
 - Keras
- ONNX

Resources

- <http://www.deeplearningbook.org/>
- https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html
- <https://www.tensorflow.org/tutorials>
- StackOverflow
- Elements of Statistical Learning, Chapter 11